

Machine Learning Tutorial – Support Vector Machine (SVM)

Module: Machine Learning and Neural Networks

Assignment Type: Individual Report

Student Name: Lohith Kamuju

Student ID: 23066840

Submission Date: 27 March 2025

Tutor: Peter Scicluna

GitHub link: https://github.com/lohithkamuju7272/SVM_Tutorial

Table of Contents

- 1. Introduction**
- 2. Theoretical Foundation**
- 3. Practical Implementation**
- 4. Kernel Comparison and Results**
- 5. Discussion and Best Practices**
- 6. Advantages, Limitations, Disadvantages and Challenges**
- 7. Future Enhancements**
- 8. Conclusion**
- 9. References**

Machine Learning Tutorial – Support Vector Machine (SVM)

1.Introduction:

Support Vector Machine (SVM) is a robust classification algorithm that finds the optimal boundary called a **hyperplane** to separate different classes. It focuses on the most critical data points, known as **support vectors**, to maximize the margin between classes, improving both accuracy and generalization.

Even for complex, non-linear data, SVM uses the **kernel trick** to map inputs into higher dimensions where separation becomes easier. This report explores SVM's core principles, kernel variations, and model tuning, supported by practical implementation.

Visualizing SVM Margins

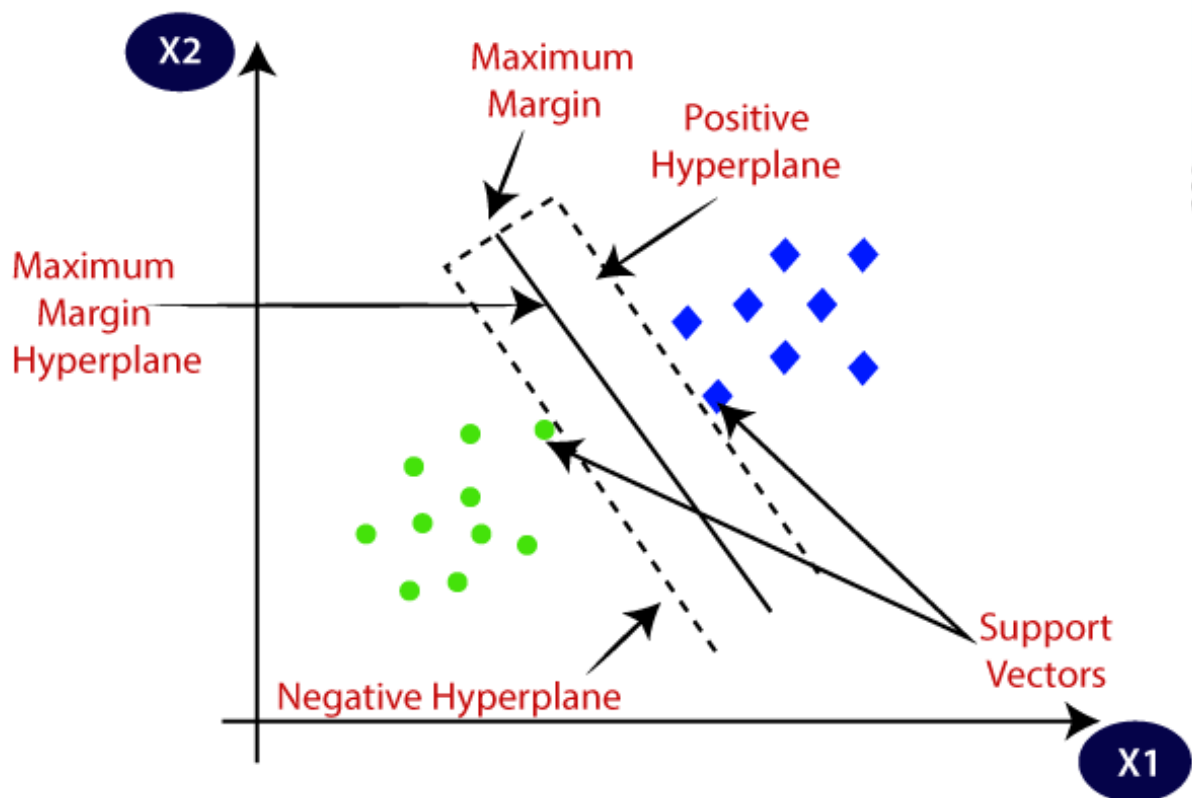


Figure 1: Illustration of maximum margin hyperplanes in SVM.
The nearest data points, known as support vectors, dictate the margin boundaries.

The image shows two classes (green and blue) separated by a central hyperplane. The dashed lines represent the maximum margin, and the closest points **support vectors** define this space. SVM uses them to build the most effective decision boundary.

Machine Learning Tutorial – Support Vector Machine (SVM)

2.Theoretical Foundation

2.1 The SVM Model

SVMs seek to maximize the margin between classes by finding an optimal hyperplane in the feature space. For linearly separable data, the formulation can be expressed as:

$$\min_{w, b} \frac{1}{2} \|w\|^2$$

subject to

$$y_i (W \cdot X_i + b) \geq 1 \quad \forall i$$

where:

- w is the weight vector defining the orientation of the hyperplane,
- b is the bias term shifting the hyperplane,
- y_i are class labels (commonly +1 or -1),
- x_i are the feature vectors.

Maximizing the margin often yields better generalization to new data

2.2 Soft Margin and Slack Variables

Most real-world data are not perfectly separable. To accommodate some misclassifications, SVMs introduce slack variables (ξ_i). The optimization problem becomes:

$$\min_{w, b} \frac{1}{2} \|w\|^2 + C \sum_{i=1}^n \xi_i$$

subject to

$$y_i (W \cdot X_i + b) \geq 1 - \xi_i, \quad \xi_i \geq 0$$

The parameter C controls the trade-off between margin width and classification errors. Higher values prioritize fewer errors, while lower values allow a wider margin with more tolerance for mistakes.

2.3 The Kernel Trick

When data are not linearly separable in their original space, the kernel trick maps them into a higher-dimensional space (implicitly) where a linear boundary may be feasible. Common kernels include:

- **Linear Kernel:**

$$K(x, x') = x \cdot x'$$

Best for data close to linearly separable.

- **Polynomial Kernel:**

$$K(x, x') = (x \cdot x' + c)^d$$

- c : Constant offset
- d : Polynomial degree
Captures polynomial-like relationships in data.

Machine Learning Tutorial – Support Vector Machine (SVM)

- **RBF (Gaussian) Kernel:**

$$K(x, x') = \exp(-\gamma \|x - x'\|^2)$$

- γ : Determines the influence of each training sample
A flexible default for modeling complex, non-linear boundaries.

2.4 Hyperparameter Tuning and Model Complexity

Performance depends on carefully tuning **CCC**, γ /**gamma** (for RBF), and the polynomial **degree** (for polynomial kernels). Using techniques like GridSearchCV helps find a balance between underfitting and overfitting.

2.5 Geometric Interpretation and Support Vectors

Only the data points closest to the decision boundary—known as support vectors—determine the margin. This characteristic makes SVMs both efficient and interpretable, as they focus on a critical subset of points rather than the entire dataset.

2.6 Extensions to Multi-Class Classification

Although SVMs are inherently binary classifiers, strategies like **one-vs-one** or **one-vs-rest** allow them to handle multi-class tasks by training multiple binary classifiers and combining their outputs.

3. Practical Implementation

While theory builds understanding, hands-on implementation reveals deeper insights. Here, we apply SVMs on a non-linear make_circles dataset from Scikit-learn to compare the performance of linear, polynomial, and RBF kernels.

Our approach includes:

- Generate and preprocess data for better performance.
- Train SVMs with various kernels for comparison.
- Tune hyperparameters via GridSearchCV for accuracy.
- Visualize decision boundaries to understand model behaviour.

```
1 # SVM Kernel Comparison.py
2 # Exploring the impact of different kernel functions in SVM for classification
3
4 import numpy as np
5 import matplotlib.pyplot as plt
6 from sklearn.datasets import make_circles
7 from sklearn.model_selection import train_test_split, GridSearchCV
8 from sklearn.preprocessing import StandardScaler
9 from sklearn.svm import SVC
10 from sklearn.metrics import classification_report, confusion_matrix
11
12 # Generate a synthetic non-linear dataset (two interleaving circles)
13 X, y = make_circles(n_samples=500, noise=0.2, factor=0.5, random_state=42)
14
15 # Split dataset into training (70%) and testing (30%) sets
16 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
17
18 # Standardize feature values to improve model performance
19 scaler = StandardScaler()
20 X_train_scaled = scaler.fit_transform(X_train)
21 X_test_scaled = scaler.transform(X_test)
22
23 # Train SVM models with different kernels
24 kernels = ['linear', 'poly', 'rbf']
25 models = {}
26
27 for kernel in kernels:
28     if kernel == 'poly':
29         model = SVC(kernel=kernel, degree=3, C=1.0, random_state=42) # Polynomial kernel
30     else:
31         model = SVC(kernel=kernel, C=1.0, random_state=42) # Linear & RBF kernels
32     model.fit(X_train_scaled, y_train)
33     models[kernel] = model
34
35 # Evaluate models
36 for kernel, model in models.items():
37     y_pred = model.predict(X_test_scaled)
38     report = classification_report(y_test, y_pred)
39     print(f'Classification Report for {kernel} kernel:\n{report}\n')
```

weighted avg 0.884 0.880 0.880 150

Best Parameters from Grid Search: {'C': 1, 'gamma': 'scale', 'kernel': 'rbf'}

Best Accuracy from Grid Search: 0.87

In [2]:

Variable	Type	Value
clf	svm_classes.SVC	1 SVC object of skl...
grid_search	model_selection_search.GridSearchCV	1 GridSearchCV obje...
kernel	str	3 rbf
kernels	list	3 ['linear', 'poly', 'rbf']
model	svm_classes.SVC	1 SVC object of skl...
models	dict	3 {'linear': SVC, 'p...
param_grid	dict	3 {'C': [0.1, 1, 10]}
scaler	preprocessing_data.StandardScaler	1 StandardScaler ob...
X	Array of float64	(500, 2) [[-0.46908421 0. ...
X_test	Array of float64	(150, 2) [[0.12386702 1. ...
X_test_scaled	Array of float64	(150, 2) [[0.40939467 0. ...
X_train	Array of float64	(350, 2) [[0.78366013 0. ...

Name	Size	Last Modified
Assignment 1 Draft.pdf	892.92 KIB	21/03/2025 23:41
Assignment 1.docx	681.04 KIB	21/03/2025 23:24
custom_svm_dataset.csv	20.33 KIB	20/03/2025 22:15
draft.docx	601.61 KIB	26/03/2025 21:47
make_circles_dataset.py	2.86 KIB	20/03/2025 22:33
SVM Tutorial.docx	445.10 KIB	26/03/2025 23:14

Machine Learning Tutorial – Support Vector Machine (SVM)

```
Spyder (Python 3.12)
File Edit Search Source Run Debug Consoles Projects Tools View Help

C:\Assignments\SVM Tutorial\make_circles_dataset.py

29 model = SVC(kernel='kernel', degree=3, C=1.0, random_state=42) # Polynomial kernel
30 else:
31     model = SVC(kernel='kernel', C=1.0, random_state=42) # Linear & RBF kernels
32
33 model.fit(X_train_scaled, y_train)
34 models[kernel] = model
35
36 # Evaluate model performance
37 for kernel, clf in models.items():
38     y_pred = clf.predict(X_test_scaled)
39     print(f" SVM with {kernel.upper()} Kernel ==")
40     print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
41     print("Classification Report:\n", classification_report(y_test, y_pred, digits=3))
42
43 # Perform hyperparameter tuning using GridSearchCV
44 param_grid = {
45     'C': [0.1, 1, 10],
46     'gamma': ['scale', 'auto'],
47     'kernel': ['linear', 'poly', 'rbf']
48 }
49
50 grid_search = GridSearchCV(SVC(), param_grid, cv=5, scoring='accuracy')
51 grid_search.fit(X_train_scaled, y_train)
52
53 print("Best Parameters from Grid Search:", grid_search.best_params_)
54 print("Best Accuracy from Grid Search: {:.3f}".format(grid_search.best_score_))
55
56 # Function to visualize decision boundaries
57 def plot_decision_boundary(model, X, y, title):
58     x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
59     y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
60     xx, yy = np.meshgrid(np.linspace(x_min, x_max, 200),
61                          np.linspace(y_min, y_max, 200))
62
63     Z = model.predict(np.c_[xx.ravel(), yy.ravel()])
64     Z = Z.reshape(xx.shape)
65
66     plt.figure(figsize=(6, 4))
67     plt.contourf(xx, yy, Z, alpha=0.3, cmap='coolwarm')
68     plt.scatter(X[:, 0], X[:, 1], c=y, edgecolors='k', cmap='coolwarm')
69     plt.title(title)
70     plt.xlabel('Feature 1')
71     plt.ylabel('Feature 2')
72     plt.show()
73
74 # Visualize decision boundaries for each kernel
75 for kernel, clf in models.items():
76     plot_decision_boundary(clf, X_test_scaled, y_test, f"SVM with {kernel.upper()} Kernel")
77
```

This code implements Support Vector Machines (SVMs) with different kernels (Linear, Polynomial, and RBF) on a non-linear dataset (*make_circles*), showcasing how each kernel affects the decision boundary and classification performance.

Figure 1: SVM with Linear Kernel

This plot shows the SVM with a linear kernel trying to split two classes using a straight line. However, the data is non-linear, leading to poor separation and frequent misclassifications—making the linear kernel a poor fit for this task.

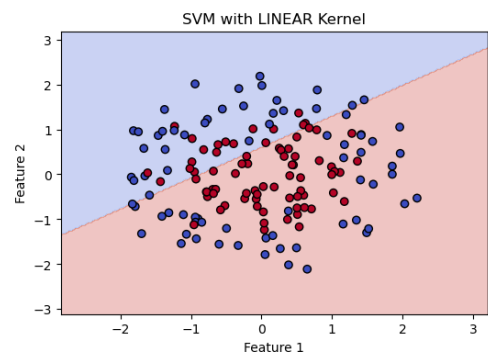


Figure 2: SVM with Polynomial Kernel

This plot shows an SVM with a polynomial kernel forming a curved boundary to capture non-linear patterns. While it improves class separation, higher polynomial degrees can make the model overly complex and prone to overfitting.

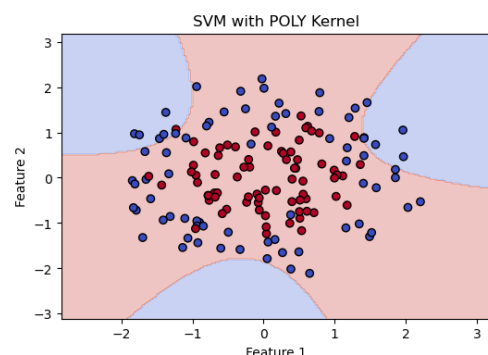



Figure 3: SVM with RBF Kernel

This plot shows an SVM with an RBF kernel creating a flexible boundary by mapping data into higher dimensions. It captures complex patterns well, offering the best separation for this non-linear dataset.



Machine Learning Tutorial – Support Vector Machine (SVM)

4. Discussion and Best Practices

Comparing SVM Kernel Performance

Kernel Comparison

Each kernel behaved differently in our implementation:

- **Linear:** Performed poorly on the non-linear dataset, causing frequent misclassifications.
- **Polynomial:** Improved separation with curved boundaries but risked overfitting at high degrees.
- **RBF:** Delivered the best results by creating flexible, well-separated decision boundaries.

Best Practices for SVM

- **Pick the Right Kernel:** Use linear for simple data, RBF or polynomial for complex patterns.
- **Scale Features:** Standardize input for better distance-based performance.
- **Tune Hyperparameters:** Adjust C, gamma, and degree via GridSearchCV to avoid over/underfitting.
- **Avoid Overfitting:** Keep models simple when possible—don't overcomplicate kernels.
- **Cross-Validate:** Use validation sets to ensure reliable generalization.
- **Reduce Dimensionality:** Apply PCA when dealing with many features for faster, cleaner models.

Discussion and Best Practices

This table highlights the **main challenges of SVMs** and practical ways to overcome them. By tackling issues like **high computation costs**, **sensitivity to hyperparameters**, **overlapping classes**, and **feature overload**, we can enhance SVM efficiency and improve classification accuracy.

Challenge	Description	Best Approach to Overcome It
Computational Intensity	<ul style="list-style-type: none">- SVMs require significant processing power for large datasets.- Training becomes slow due to complex mathematical optimization.	<ul style="list-style-type: none">- Use Stochastic Gradient Descent (SGD) to speed up training.- For massive datasets, switch to Random Forest or Neural Networks for better scalability.
Hyperparameter Sensitivity	<ul style="list-style-type: none">- Model performance heavily depends on C, gamma, and kernel selection.- Poor tuning can lead to underfitting or overfitting.	<ul style="list-style-type: none">- Use GridSearchCV or RandomizedSearchCV for systematic tuning.- Apply cross-validation to select the best hyperparameters.
Struggles with Overlapping Classes	<ul style="list-style-type: none">- SVM performs best when there's a clear margin between classes.- If classes are mixed, even non-linear kernels may fail to define an optimal boundary.	<ul style="list-style-type: none">- Improve class separation using feature engineering.- Use ensemble learning techniques (e.g., boosting, bagging) to enhance classification accuracy.
Handling High-Dimensional Data	<ul style="list-style-type: none">- When the number of features is greater than the number of samples, SVMs can easily overfit.- Too many irrelevant features reduce model efficiency.	<ul style="list-style-type: none">- Apply dimensionality reduction techniques like PCA (Principal Component Analysis).- Use feature selection methods to keep only the most important variables.

Machine Learning Tutorial – Support Vector Machine (SVM)

5. Advantages, Limitation, Disadvantages, and Challenges of SVM

This table outlines the key strengths and challenges of SVM, showcasing its ability to handle complex, high-dimensional data while also addressing its computational demands and need for careful tuning. Choosing the right kernel and hyperparameters is essential to unlocking SVM's full potential.

Category	Key Points
Advantages	<ul style="list-style-type: none">▪ Effective in high-dimensional spaces like text and image data.▪ Kernel trick enables handling of non-linear patterns.▪ Maximizes margin to reduce overfitting and boost generalization.▪ Performs well on small to medium datasets.▪ Supports various kernels for flexible modelling.
Limitations	<ul style="list-style-type: none">▪ Computationally intensive on large datasets due to complex calculations.▪ Requires careful tuning; poor settings may cause under/overfitting.▪ Struggles with overlapping classes where boundaries are unclear.▪ Sensitive to noise and outliers affecting boundary placement.▪ Less interpretable compared to decision trees with clear logic.
Disadvantages	<ul style="list-style-type: none">▪ Training is slow on large data due to quadratic optimization.▪ Not suited for real-time tasks needing fast predictions.▪ Needs feature scaling (e.g., StandardScaler) for best results.▪ Prone to overfitting with too many irrelevant features.▪ Models like Random Forest, XGBoost, or Neural Nets often scale better.
Challenges	<ul style="list-style-type: none">▪ Kernel choice is crucial; wrong picks hurt accuracy.▪ Training slows with larger datasets.▪ Struggles with imbalanced data, often favouring the majority class.▪ High-degree polynomials risk overfitting on small data.▪ Less interpretable than models like decision trees.

Machine Learning Tutorial – Support Vector Machine (SVM)

6.Future Enhancements

As Machine learning continues to evolve, it's essential for SVM to adapt accordingly. The following forward-looking improvements aim to enhance SVM's performance, make it more scalable, and expand its application across a wide range of real-world problems.

A. Feature Engineering for Better Performance

- Extracting more meaningful features can significantly enhance classification accuracy.
- Polynomial features or domain-specific transformations can provide better input representations.

B. Hybrid Approaches

- Combining SVMs with Neural Networks can boost performance in areas like image recognition, speech processing, and sentiment analysis.
- Hybrid models can leverage SVM's ability to find optimal margins and Neural Networks' ability to extract deep features.

C. Custom Kernel Functions

- For specialized problems in finance, medicine, and bioinformatics, designing custom kernels tailored to domain-specific data structures improves classification accuracy.

D. Efficient Training Techniques

- Reducing computational cost using Stochastic Gradient Descent (SGD) or Approximate SVMs can make training feasible for large-scale problems.
- Implementing parallel processing techniques speeds up training in big data environments.

E. Integration with Big Data Tools

- Adapting SVMs to work with distributed systems like Apache Spark or Dask makes them scalable to enterprise-level problems.

F. Exploring Alternative Models

- If the dataset is extremely large, models like Gradient Boosting, Decision Trees, or Deep Learning may provide better scalability and performance.

By embracing these improvements, SVMs can continue to thrive in data-rich environments, ensuring they remain a reliable and adaptable choice for solving modern classification challenges.

7.Conclusion

Support Vector Machines have proven to be a powerful classification tool, particularly when handling complex, non-linear datasets. This tutorial highlighted their practical application, theoretical foundation, and performance under different kernel settings, with the RBF kernel delivering the best results for the selected dataset.

Key takeaways from this tutorial include:

1. **Adaptability Through Kernels:** SVMs can tackle a wide range of problems by leveraging various kernel functions that enable linear separation in higher dimensions.
2. **Effectiveness in Non-Linear Scenarios:** The RBF and polynomial kernels outperform linear ones in cases where the data structure is curved or intertwined.
3. **Importance of Hyperparameter Tuning:** Properly adjusting parameters like C and gamma is crucial for achieving optimal performance and avoiding overfitting or underfitting.
4. **Visualization Aids Understanding:** Graphical representation of decision boundaries offers valuable insights into how different models behave with the same dataset.
5. **Room for Advancement:** Despite its strengths, SVM is not without limitations. Incorporating modern techniques like feature engineering, hybrid models, and big data integration can unlock even greater potential.

In conclusion, SVMs remain a relevant and valuable method in the data science toolkit. With thoughtful implementation and continuous improvement, they can deliver highly accurate results across a wide variety of domains.

Machine Learning Tutorial – Support Vector Machine (SVM)

8. References

The following sources were referenced for theoretical concepts, practical implementation, datasets, and visuals used in this report.

Academic References

- Bishop, C. M. (2006). *Pattern Recognition and Machine Learning*. Springer.
- Cortes, C., & Vapnik, V. (1995). *Support-vector networks*. *Machine Learning*, 20(3), 273–297. <https://doi.org/10.1007/BF00994018>
- Hastie, T., Tibshirani, R., & Friedman, J. (2009). *The Elements of Statistical Learning: Data Mining, Inference, and Prediction* (2nd ed.). Springer.
- Jolliffe, I. T. (2002). *Principal Component Analysis* (2nd ed.). Springer Series in Statistics.
- Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., ... & Duchesnay, É. (2011). *Scikit-learn: Machine Learning in Python*. *Journal of Machine Learning Research*, 12, 2825–2830.

Online Resources

- Scikit-learn Documentation. Available at: <https://scikit-learn.org>
- Scikit-learn. (n.d.). *make_circles* - Generate a large circle containing a smaller circle in two-dimensional space.
Available at: https://scikit-learn.org/stable/modules/generated/sklearn.datasets.make_circles.html
- StatQuest with Josh Starmer (YouTube). Available at: <https://www.youtube.com/user/joshstarmer>
- Coursera – Machine Learning by Andrew Ng. Available at: <https://www.coursera.org/learn/machine-learning>

Web Articles

- Deshpande, V. (2021). *Understanding the Inner Workings of SVM: From Data Preprocessing to Model Deployment*. Medium. Available at: <https://medium.com/@vdeshpande551/understanding-the-inner-workings-of-svm-from-data-preprocessing-to-model-deployment-cdc9d72a2d34>

Figures

- Figure 1 in the Introduction adapted from: Deshpande, V. (2021). *Understanding the Inner Workings of SVM*. Medium. <https://medium.com/@vdeshpande551/understanding-the-inner-workings-of-svm-from-data-preprocessing-to-model-deployment-cdc9d72a2d34>